

Homework 3: Part II

Louis Bensard

February 27, 2018

(d)

```
p = 3
n = 200
m = p*(p+1)/2
maxit = 1000

mu = c(-1,1,2)
Sigma = matrix(c(1,0.7,0.7,0.7,1,0.7,0.7,0.7,1),3)
theta = vectorize(mu, Sigma,p)

#data
x = gen(n, p, mu, Sigma)

mu_0 = c(-1.5,1.5,2.3)
Sigma_0 = matrix(c(1,0.5,0.5,0.5,1,0.5,0.5,0.5,1),3)
theta_0 = vectorize(mu_0, Sigma_0, p)

tolgrad = 1e-9 ; tolerr = 1e-6
```

I - Steepest Ascent method with step halving

```
max_seeker(x, mu_0, Sigma_0, n, maxit, TRUE, FALSE, FALSE, tolerr, tolgrad)
```

```
##
## Iteration    Halving    log-likelihood    ||gradient||
##      1              -853.5459      1.0e+03
##      1          1             NA         NA
##      1          2             NA         NA
##      1          3             NA         NA
##      1          4             NA         NA
##      1          5             NA         NA
##      1          6             NA         NA
##      1          7             NA         NA
##      1          8             NA         NA
##      1          9             NA         NA
## -----
##      2          -783.9008      1.5e+03
##      2          1             NA         NA
##      2          2             NA         NA
##      2          3             NA         NA
##      2          4             NA         NA
##      2          5             NA         NA
##      2          6             NA         NA
##      2          7             NA         NA
```

##	2	8	NA	NA
##	2	9	-850.7135	1.2e+02
##	2	10	-808.6329	1.3e+02
##	2	11	-777.6896	1.5e+02
##	-----			
##	...			
##	...			
##	...			
##	998		-709.9636	1.7e-05
##	998	1	-709.9636	1.0e-02
##	998	2	-709.9636	5.2e-03
##	998	3	-709.9636	2.6e-03
##	998	4	-709.9636	1.3e-03
##	998	5	-709.9636	6.4e-04
##	998	6	-709.9636	3.2e-04
##	998	7	-709.9636	1.5e-04
##	998	8	-709.9636	7.2e-05
##	998	9	-709.9636	3.1e-05
##	998	10	-709.9636	1.1e-05
##	998	11	-709.9636	8.6e-07
##	998	12	-709.9636	4.2e-06
##	998	13	-709.9636	6.8e-06
##	998	14	-709.9636	8.1e-06
##	998	15	-709.9636	8.7e-06
##	998	16	-709.9636	9.0e-06
##	-----			
##	999		-709.9636	1.6e-05
##	999	1	-709.9636	1.0e-02
##	999	2	-709.9636	5.0e-03
##	999	3	-709.9636	2.5e-03
##	999	4	-709.9636	1.3e-03
##	999	5	-709.9636	6.2e-04
##	999	6	-709.9636	3.1e-04
##	999	7	-709.9636	1.5e-04
##	999	8	-709.9636	7.0e-05
##	999	9	-709.9636	3.0e-05
##	999	10	-709.9636	1.1e-05
##	999	11	-709.9636	8.3e-07
##	999	12	-709.9636	4.1e-06
##	999	13	-709.9636	6.6e-06
##	999	14	-709.9636	7.8e-06
##	999	15	-709.9636	8.4e-06
##	999	16	-709.9636	8.7e-06
##	999	17	-709.9636	8.9e-06
##	999	18	-709.9636	8.9e-06
##	999	19	-709.9636	9.0e-06
##	-----			
##	1000		-709.9636	1.6e-05
##	1000	1	-709.9636	1.0e-02
##	1000	2	-709.9636	5.0e-03
##	1000	3	-709.9636	2.5e-03
##	1000	4	-709.9636	1.2e-03
##	1000	5	-709.9636	6.2e-04
##	1000	6	-709.9636	3.1e-04

```

## 1000      7      -709.9636      1.5e-04
## 1000      8      -709.9636      7.0e-05
## 1000      9      -709.9636      3.0e-05
## 1000     10      -709.9636      1.1e-05
## 1000     11      -709.9636      8.3e-07
## 1000     12      -709.9636      4.1e-06
## 1000     13      -709.9636      6.5e-06
## 1000     14      -709.9636      7.8e-06
## 1000     15      -709.9636      8.4e-06
## 1000     16      -709.9636      8.7e-06
## 1000     17      -709.9636      8.8e-06
## 1000     18      -709.9636      8.9e-06
## -----
## -----

```

After 1000 iterations, the condition on the norm of the gradient ($<1e-9$) hasn't been met, therefore the algorithm stops at $\text{maxit} = 1000$. If we would increase tolgrad to say $1e-5$, the algorithm would stop after 323 iterations. But for our case, the maximum -709.9636 of the log-likelihood function is reached after 1000 iterations.

II - Fisher-scoring method with step-halving

```
max_seeker(x, mu_0, Sigma_0, n, maxit, FALSE, TRUE, FALSE, tolerr, tolgrad)
```

```

##
## Iteration   Halving   log-likelihood   ||gradient||
##      1           0      -853.5459      1.0e+03
##      1           0      -745.4340      8.3e+01
## -----
##      2           0      -745.4340      1.6e+02
##      2           0      -709.9636      1.2e-12
## -----
## ...
## ...
## ...
##      1           0      -853.5459      1.0e+03
##      1           0      -745.4340      8.3e+01
## -----
##      2           0      -745.4340      1.6e+02
##      2           0      -709.9636      1.2e-12
## -----
##      3           0      -709.9636      2.5e-12
##      3           0      -709.9636      2.0e-13
## -----
## -----

```

Both the norm of the gradient and the Modified Relative Error meet the criteria to stop the algorithm after 3 iterations. The maximum -709.9636 of the log-likelihood function is reached after 3 iterations.

III - Newton's method with step-halving

```
max_seeker(x, mu_0, Sigma_0, n, maxit, FALSE, FALSE, TRUE, tolerr, tolgrad)
```

```
##
## Iteration    Halving    log-likelihood    ||gradient||
##      1              -853.5459      1.0e+03
##      1          1          NA          NA
##      1          2          NA          NA
##      1          3          NA          NA
##      1          4          NA          NA
## -----
##      2              -848.6825      4.4e+03
##      2          1          NA          NA
## -----
## ...
## ...
## ...
##      7              -709.9694      1.1e+01
##      7          0      -709.9636      9.9e-02
## -----
##      8              -709.9636      2.1e-01
##      8          0      -709.9636      4.3e-05
## -----
##      9              -709.9636      9.3e-05
##      9          0      -709.9636      6.9e-12
## -----
## -----
```

Both the norm of the gradient and the Modified Relative Error meet the criteria to stop the algorithm after 9 iterations. The maximum -709.9636 of the log-likelihood function is reached after 9 iterations.

(e)

```
mu = c(-1,1,2)
Sigma_e = matrix(c(1,0.9,0.9,0.9,1,0.9,0.9,0.9,1),3)

#data
x_e = gen(n, p, mu, Sigma_e)

mu_0 = c(-1.5,1.5,2.3)
Sigma_0 = matrix(c(1,0.5,0.5,0.5,1,0.5,0.5,0.5,1),3)
theta_0 = vectorize(mu_0, Sigma_0, p)

max_seeker(x_e, mu_0, Sigma_0, n, maxit, TRUE, FALSE, FALSE, tolerr, tolgrad)
```

```
##
## Iteration    Halving    log-likelihood    ||gradient||
##      1              -791.1402      9.4e+02
##      1          1          NA          NA
##      1          2          NA          NA
##      1          3          NA          NA
##      1          4          NA          NA
##      1          5          NA          NA
##      1          6          NA          NA
##      1          7          NA          NA
##      1          8          NA          NA
```

##	1	9	NA	NA
##	1	10	NA	NA
##	-----			
##	2		-670.8129	9.4e+02
##	2	1	NA	NA
##	2	2	NA	NA
##	2	3	NA	NA
##	2	4	NA	NA
##	2	5	NA	NA
##	2	6	NA	NA
##	2	7	NA	NA
##	2	8	NA	NA
##	2	9	NA	NA
##	2	10	NA	NA
##	2	11	NA	NA
##	-----			
##	...			
##	...			
##	...			
##	998		-505.6665	5.0e+00
##	998	1	NA	NA
##	998	2	NA	NA
##	998	3	NA	NA
##	998	4	-564.1572	1.0e+03
##	998	5	-529.0145	4.7e+02
##	998	6	-514.1718	3.5e+02
##	998	7	-508.3743	2.5e+02
##	998	8	-506.4386	1.5e+02
##	998	9	-505.8697	8.3e+01
##	998	10	-505.7164	4.3e+01
##	998	11	-505.6777	2.1e+01
##	998	12	-505.6685	9.5e+00
##	998	13	-505.6666	3.6e+00
##	998	14	-505.6663	8.1e-01
##	-----			
##	999		-505.6663	1.4e+00
##	999	1	NA	NA
##	999	2	NA	NA
##	999	3	NA	NA
##	999	4	-535.1274	2.7e+03
##	999	5	-510.1268	6.1e+02
##	999	6	-506.5793	2.3e+02
##	999	7	-505.8728	9.9e+01
##	999	8	-505.7145	4.6e+01
##	999	9	-505.6775	2.2e+01
##	999	10	-505.6687	1.1e+01
##	999	11	-505.6668	5.0e+00
##	999	12	-505.6663	2.3e+00
##	999	13	-505.6663	1.0e+00
##	-----			
##	1000		-505.6663	2.3e+00
##	1000	1	NA	NA
##	1000	2	-594.1175	1.9e+03
##	1000	3	-539.8014	5.8e+02

```
## 1000      4      -518.7646      3.9e+02
## 1000      5      -510.0778      2.9e+02
## 1000      6      -506.9839      1.9e+02
## 1000      7      -506.0276      1.1e+02
## 1000      8      -505.7598      5.8e+01
## 1000      9      -505.6894      3.0e+01
## 1000     10      -505.6716      1.5e+01
## 1000     11      -505.6674      7.2e+00
## 1000     12      -505.6664      3.2e+00
## 1000     13      -505.6663      1.3e+00
## -----
## -----
```

After 1000 iterations, the condition on the norm of the gradient ($<1e-9$) hasn't been met, therefore the algorithm stops at `maxit = 1000`. The maximum -505.6663 of the log-likelihood function is reached after 1000 iterations.

But after 1000 iterations, in part (d) the norm of the gradient was $8.9e-06$ as opposed to 1.3 for this part (e) so we can argue that part (d) is faster than part (e). The following proves more formally that this is correct:

```
theta_e = vectorize(mu, Sigma_e,p)

ev1 = eigen(-hessian(x, theta))[[1]]
ev2 = eigen(-hessian(x_e, theta_e))[[1]]

cr1 = ((max(ev1) - min(ev1))/(max(ev1) + min(ev1)))^2
cr2 = ((max(ev2) - min(ev2))/(max(ev2) + min(ev2)))^2

cat("cr1 =", cr1, "\ncr2 =", cr2)

## cr1 = 0.9508285
## cr2 = 0.9958086
```

The convergence ratio of part (d) is 0.9508, which is less than the convergence ratio of 0.9952 in part (e). Therefore, part (e) is indeed slower than part (d).

(f)

```
vect_ste = c()
mean_vect = c(mean(x[,1]), mean(x[,2]), mean(x[,3]))
Sig_cov = ((n-1)/n)*cov(x)
theta_mle = vectorize(mean_vect, Sig_cov,p)
f = fisher(x, theta_mle)
f_inv = solve(f)

for(i in 1:(m+p)){

  for(j in 1:(m+p)){

    if(i==j) vect_ste = c(vect_ste, sqrt(f_inv[i,j]))
  }
}

unvectorize(vect_ste, p)
```

```
## $mu
## [1] 0.07197462 0.07485894 0.07365860
##
## $Sigma
##           [,1]      [,2]      [,3]
## [1,] 0.10360692 0.09276979 0.09079161
## [2,] 0.09276979 0.11207721 0.09807797
## [3,] 0.09079161 0.09807797 0.10851179
```

\$mu is the standard error of the MLE $\hat{\mu}$ of μ . \$Sigma is the standard error of the MLE $\hat{\Sigma}$ of Σ .